

Robot Intelligence

Fall 2022

Midterm Exam

8/31/2022

Take Home Exam - Due Before Class on Canvas 10/28/2022

Name: _____

This exam contains 11 pages (including this cover page) and 9 questions. The total number of points is 180. Graduate students must answer some additional questions. Undergraduates answering graduate questions will be given additional points and a feeling of satisfaction from attempting difficult things (probably).

Any programming/code artifacts associated with this exam should be submitted along with the final PDF to canvas. Please put your entire submission in a single .zip file.

This exam will cover topics in motion, planning, sensor processing, and ethics.

Grade Table (Quick view to see the breakdown)

Question	Points	Score
1	20	
2	20	
3	20	
4	20	
5	20	
6	20	
7	20	
8	20	
9	20	
Total:	180	

1. (20 points) Moving in a car

I would like you to use a skid steer model of a robot that is 75cm long and 55cm wide and run a few simple experiments with it. Please upload your resulting figures and python (or other) code:

- (a) (5 points) Make a list of commands (at $t=0.1$) that will allow this robot to traverse along the edge of a 5m diameter circle. The robot starts off in the center of the circle $(0,0)$, and you cannot leave the circle's border. Plot both the resulting path (x, y) and trajectory $(x, y, \text{ and angular velocities})$. Assume a constant velocity of 8m/s
- (b) (5 points) Do the same as the above for a traditional car (Ackermann steering).
- (c) (10 points) Your Ackermann vehicle (with the same dimensions as described for the skid steer), is driving on a circle of radius 2.5m . Assume that you begin on the edge of the circle. Calculate the positional error with our computational approximation using the forward Euler method, as referred to in the course notes and illustrated in Reading 2, eq. 1.6. Graph the errors and computing time for three different time-steps ($\Delta t = 1, 0.1, 0.01$), error is defined as the absolute distance between the expected (from equations) to real x, y position (defined analytically) over time.

Brief notes on what this problem is about: Imagine that you are moving a semi-truck and you have GPS positions (somewhat accurate) and an internal prediction model. We want the internal models to be updated based on "ground truth" information to build better estimates of the vehicle motion over time.

- (d) (10 points) **Graduate Student Question** Compute part c again where road frictions are much lower (due to rain or ice). Slip causes your tires to respond very differently. Assume that your theta is $\theta_{actual} = \theta(1 - 0.08)$ and velocity is $v_{actual} = v(1 - 0.04)$.




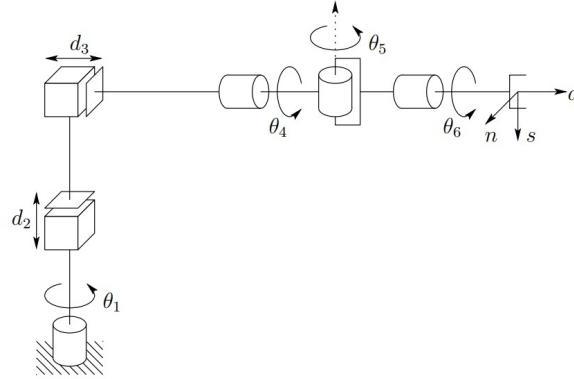
Figure 1: SCARA type manipulator.

2. (20 points) Inverse Kinematics, numerical approaches

The following SCARA-type manipulator has physical characteristics as follows: $l_1 = 60cm$, $l_2 = 40cm$, and an arbitrary height (ie: assume that $(0,0)$ is at the first motor).

- (5 points) What is the workspace volume for this robot?. (Draw a picture of it as well)
- (5 points) Write the DH parameters
- (10 points) Compute the forward kinematics to find the position of the end effector if $\theta_1 = 30deg$, $\theta_2 = 45deg$, $\theta_3 = 90deg$, and $d = 14cm$

3. (20 points) Inverse Kinematics with numerical approaches (Use the manipulator on the following page). Note that d_6 refers to the length of the arm on the top and d_1 is the length of the system from the base to the first motor.)
- (a) (10 points) Compute the joint angles and extension distances that will get this robot to pick up an object at $x=1.2$, $y=0.8$, $z=0.5$ if the robot started at if the robot started with the end effector at $\theta_1 = -90deg$, $d_2 = 0.5m$, $d_3 = 1.0m$, $\theta_4 = -90deg$, $\theta_5 = 90deg$, $\theta_6 = 40deg$, $d_6 = 0.2m$. Note that d_6 is a constant value and that θ_6 will not impact the final solution for our purposes, you are allowed to set d_1 as it only impacts the final z position.
 - (b) (10 points) What would be the joint angles and extension distances to get to the same goal coordinates if the robot started with the end effector at $\theta_1 = 0deg$, $d_2 = 0.2m$, $d_3 = 0.3m$, $\theta_4 = -90deg$, $\theta_5 = 90deg$, $\theta_6 = 40deg$, $d_6 = 0.2m$ and we wished to minimize the total distance traveled for each actuating part?
 - (c) (10 points) **Graduate Question:** Using the same system, start configuration, and goal coordinates of Part B, compute the energy-efficient joint angles and extension distances. For your energy model, assume that actuating θ_1 costs 3x the energy of $\theta_{4,5,6}$, and $d_{2,3}$ costs 2x the energy.



$$\begin{aligned}
 T_6^0 &= \begin{bmatrix} c_1 & 0 & -s_1 & -s_1 d_1 \\ s_1 & 0 & c_1 & c_1 d_3 \\ 0 & -1 & 0 & d_1 + d_2 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} c_4 c_5 c_6 - s_4 s_6 & -c_4 c_5 s_6 - s_4 c_6 & c_4 s_5 & c_4 s_5 d_6 \\ s_4 c_5 c_6 + c_4 s_6 & -s_4 c_5 s_6 + c_4 c_6 & s_4 s_5 & s_4 s_5 d_6 \\ -s_5 c_6 & s_5 c_6 & c_5 & c_5 d_6 \\ 0 & 0 & 0 & 1 \end{bmatrix} \\
 &= \begin{bmatrix} r_{11} & r_{12} & r_{13} & d_x \\ r_{21} & r_{22} & r_{23} & d_y \\ r_{31} & r_{32} & r_{33} & d_z \\ 0 & 0 & 0 & 1 \end{bmatrix}
 \end{aligned}$$

$$\begin{aligned}
 r_{11} &= c_1 c_4 c_5 c_6 - c_1 s_4 s_6 + s_1 s_5 c_6 \\
 r_{21} &= s_1 c_4 c_5 c_6 - s_1 s_4 s_6 - c_1 s_5 c_6 \\
 r_{31} &= -s_4 c_5 c_6 - c_4 s_6 \\
 r_{12} &= -c_1 c_4 c_5 s_6 - c_1 s_4 c_6 - s_1 s_5 c_6 \\
 r_{22} &= -s_1 c_4 c_5 s_6 - s_1 s_4 c_6 + c_1 s_5 c_6 \\
 r_{32} &= s_4 c_5 c_6 - c_4 c_6 \\
 r_{13} &= c_1 c_4 s_5 - s_1 c_5 \\
 r_{23} &= s_1 c_4 s_5 + c_1 c_5 \\
 r_{33} &= -s_4 s_5 \\
 d_x &= c_1 c_4 s_5 d_6 - s_1 c_5 d_6 - s_1 d_3 \\
 d_y &= s_1 c_4 s_5 d_6 + c_1 c_5 d_6 + c_1 d_3 \\
 d_z &= -s_4 s_5 d_6 + d_1 + d_2.
 \end{aligned}$$

Figure 2: Cylindrical/Prismatic manipulator. Use for problem 6.




Figure 3: Assume that the mass $m = 0.2kg$, $l = 1.0$, $M = 4.0kg$

4. (20 points) Balancing a Pole

A cart and pole suddenly appear before you, and you feel compelled to answer burning questions you have always held about these systems.

- (a) (5 points) Describe the equations of motion that govern this system
- (b) (10 points) Generate code for a controller that would be able to keep the pole balanced in the air.
- (c) (5 points) What is the maximum angle that my pole can fall to before it cannot recover if max Force $F = 6N$?

5. (20 points) Control and Reinforcement Learning

- (a) (5 points) Explain three merits and three demerits of using reinforcement learning for mechatronic systems.
- (b) (5 points) Draw a diagram for reinforcement learning and controls and contrast the two
- (c) (10 points) Use OpenAI Gym to create a trained reinforcement learning model for the cart and pole problem (CartPole-v1).

https://www.gymnasium.ml/environments/classic_control/cart_pole/

This tutorial will help you get started with using the OpenAI gym python library.
<https://blog.paperspace.com/getting-started-with-openai-gym/>

- (d) (10 points) **Graduate Question:** Implement this example of a 2D running robot (cheetah) and adapt the code to penalize hip motor movements faster than $\pi \frac{rad}{s}$.
https://github.com/openai/gym/blob/master/gym/envs/mujoco/half_cheetah.py A Virtual Machine that was made with VMWare is available on Canvas for you to use for this question if you have trouble installing Mujoco.

6. (20 points) Human Emotions

- (a) (10 points) Using the FER library in python, example in the following link: <https://towardsdatascience.com/the-ultimate-guide-to-emotion-recognition-from-facial-e> build a system that classifies human emotions and validate on video 1 and 2 from this repository: <https://github.com/rjrahul24/ai-with-python-series/tree/main/07.%20Emotion%20Recognition%20using%20Live%20Video> Generate plots of predicted emotions over time for both videos.
- (b) (5 points) Do the same as the above with a video feed from your webcam. Set your software up to allow video feed or a pre-recorded video. (In essence, make faces at yourself and make sure that your service works). Submit a short faces-recording along with the script that can read live webcam streams.
- (c) (5 points) What are the logical applications of this tool for an autonomous robot? What are the ethical and legal consequences of fielding a system that makes decisions based on this tool?
- (d) (10 points) **Graduate Question:** Set up a service that allows 2 or more faces to be processed at once.

7. (20 points) Motion Planning

- (a) (20 points) Implement an A* based planner (an example is found here: <https://medium.com/@nicholas.w.swift/easy-a-star-pathfinding-7e6689c7f7b2>) and compare it's results with the Dijkstra, A-Star, RRT-Star, Bi-Directional A-Star, and the Breadth First Search Planners from this github. <https://github.com/AtsushiSakai/PythonRobotics/tree/master/PathPlanning>. Create a table comparing the average cost of the path found over 10 iterations along with the time to convergence.

After you have implemented your planner and compared it answer the below questions.

- Which planner provided a path with the lowest cost on average?
 - Which one found a path the fastest on average?
 - After comparing your planner to these five other ones is there anything you would change in your planner to help it converge faster or find a path with a better cost?
 - Which planner appears to be the best overall? Which planner would you use for a robot in a complex environment?
- (b) (10 points) **Graduate Question:** Implement a skid-steer (4 wheel) vehicle model for a robot 33 inches wide, by 40 inches long for the Hybrid A* algorithm. Change the scenario so both your A* and Hybrid A* have the same obstacle field. Compare and contrast the final trajectories, graph these. Write what the algorithmic differences are between the two planners and how the kinematics are used to constrain the Hybrid A*.

8. (20 points) Object Detection (Please put the classified images as part of your submission)
- (a) (5 points) Classify the first 10 pictures in https://github.com/ravirajsinh45/Crop_and_weed_detection using any image classification algorithm.
 - (b) (5 points) Implement Yolo and do the same, what are the differences.
 - (c) (10 points) Using transfer learning, pick an image classification algorithm and re-train it to learn to detect a new object of interest.
 - (d) (10 points) **Graduate Question:** Implement Detectron (<https://github.com/facebookresearch/Detectron>) and test the image segmentation of the objects you are interested in. Give metrics on relative segmentation/classification quality comparing: Mask R-CNN, faster R-CNN, and RetinaNet.

9. (20 points) Ethics of Robotics - Open Ended Questions

- (a) (5 points) Many people (particularly those in the robotics industry) believe that robotics is purely within the purview of technical development and should not have any ethical considerations. What do you feel can be a merit or demerit to this way of thinking?
- (b) (5 points) Isaac Asimov listed 3 laws of robotics, comment on the algorithmic complexity of implementing these into working intelligence. Define a scenario and write Psuedo code to implement these rules.
- (c) (5 points) In the event of an autonomous system causing harm or damages, who is responsible? Read and comment on the following two documents: <https://www.callahan-law.com/articles-and-expert-advice/when-an-autonomous-vehicle-hits-a-pe> and https://en.wikipedia.org/wiki/Self-driving_car_liability
- (d) (5 points) What laws may be helpful for regulating or controlling autonomous systems? What drawbacks will this potentially have?